

MEMORY CORE

Multi-agent dev teams burn tokens passing context through prompts.

And lose it every session reset.

- Built solo by Evan, 16 — out of China.

LIVE DEMO

Three agents. One memory layer.

planner

writes: plan
mimo-v2.5

coder

writes: file artifacts
Qwen2.5-Coder-32B

reviewer

writes: review
mimo-v2.5

→ memory-core.chinasourcingbridge.com

WHY THIS WORKS

**Role-aware memory.
Not prompts.**

LONGMEMEVAL-S

99.4

R@10 · SOTA · verified eval

LOCOMO

88.8

R@10 · published benchmark

10+ adapters · pluggable backends · Apache-2.0 eval suite · production-shaped SDK

BUILT FOR THE FEATHERLESS TRACK

Domain-specialized. Async-first. Permissive. Production-shaped.

✓ **Domain**

Software dev teams — planner, coder, reviewer.

✓ **Permissive**

MIT core + Apache-2.0 eval. Real OSS, not source-available.

✓ **Async**

Agents communicate through namespaces, not chat turns.

✓ **Production-shaped**

REST API, SDK, 10+ backends, role-scoped namespaces.

Also Collaborative Systems · Agentic Workflows.

MEMORY CORE

The memory layer for agent teams.

repo github.com/Evanyuan-builder/memory-core

demo memory-core.chinasourcingbridge.com

Built solo by Evan, 16 — out of China. Open source, MIT.